

...Then, let's load the needed packages:

```
library(tidyverse)
library(data.table)
library(tidytable)
library(readr)
```

and let's take a look at the data a little bit:

```
energy <- read.csv("~/Downloads/energydata_complete.csv")
```

```
head(energy)
```

```
##           date Appliances lights    T1    RH_1    T2    RH_2    T3
## 1 2016-01-11 17:00:00         60     30 19.89 47.59667 19.2 44.79000 19.79
## 2 2016-01-11 17:10:00         60     30 19.89 46.69333 19.2 44.72250 19.79
## 3 2016-01-11 17:20:00         50     30 19.89 46.30000 19.2 44.62667 19.79
## 4 2016-01-11 17:30:00         50     40 19.89 46.06667 19.2 44.59000 19.79
## 5 2016-01-11 17:40:00         60     40 19.89 46.33333 19.2 44.53000 19.79
## 6 2016-01-11 17:50:00         50     40 19.89 46.02667 19.2 44.50000 19.79
##           RH_3    T4    RH_4    T5    RH_5    T6    RH_6    T7
RH_7
## 1 44.73000 19.00000 45.56667 17.16667 55.20 7.026667 84.25667 17.20000
41.62667
## 2 44.79000 19.00000 45.99250 17.16667 55.20 6.833333 84.06333 17.20000
41.56000
## 3 44.93333 18.92667 45.89000 17.16667 55.09 6.560000 83.15667 17.20000
41.43333
## 4 45.00000 18.89000 45.72333 17.16667 55.09 6.433333 83.42333 17.13333
41.29000
## 5 45.00000 18.89000 45.53000 17.20000 55.09 6.366667 84.89333 17.20000
41.23000
## 6 44.93333 18.89000 45.73000 17.13333 55.03 6.300000 85.76667 17.13333
41.26000
##           T8    RH_8    T9    RH_9    T_out Press_mm_hg RH_out Windspeed
Visibility
## 1 18.2 48.90000 17.03333 45.53 6.600000      733.5      92 7.000000
63.00000
## 2 18.2 48.86333 17.06667 45.56 6.483333      733.6      92 6.666667
59.16667
## 3 18.2 48.73000 17.00000 45.50 6.366667      733.7      92 6.333333
55.33333
## 4 18.1 48.59000 17.00000 45.40 6.250000      733.8      92 6.000000
51.50000
## 5 18.1 48.59000 17.00000 45.40 6.133333      733.9      92 5.666667
47.66667
## 6 18.1 48.59000 17.00000 45.29 6.016667      734.0      92 5.333333
43.83333
##    Tdewpoint    rv1    rv2
## 1         5.3 13.27543 13.27543
## 2         5.2 18.60619 18.60619
## 3         5.1 28.64267 28.64267
## 4         5.0 45.41039 45.41039
## 5         4.9 10.08410 10.08410
## 6         4.8 44.91948 44.91948
```

As you can see, this data is wide, and not long. Variables, or features, T1 to T9 provide the

temperature of 9 rooms, and RH_1 to RH_9 provide the humidity of the same 9 rooms.

What if I'd like to make a plot of each room's temperature throughout the year? In this format, it is not possible. So let's reshape this a little bit:

```
flat_energy <- energy %>%
  pivot_longer(cols = matches("T\\d{1}"), names_to = "temperature", values_to =
"temp_value") %>%
  pivot_longer(cols = matches("RH_\\d{1}"), names_to = "humidity", values_to =
"hum_value") %>%
  mutate(temperature = case_when(temperature == "T1" ~ "kitchen",
                                temperature == "T2" ~ "living",
                                temperature == "T3" ~ "laundry",
                                temperature == "T4" ~ "office",
                                temperature == "T5" ~ "bathroom",
                                temperature == "T6" ~ "north",
                                temperature == "T7" ~ "ironing",
                                temperature == "T8" ~ "teenager",
                                temperature == "T9" ~ "parents")) %>%
  mutate(humidity = case_when(humidity == "RH_1" ~ "kitchen",
                              humidity == "RH_2" ~ "living",
                              humidity == "RH_3" ~ "laundry",
                              humidity == "RH_4" ~ "office",
                              humidity == "RH_5" ~ "bathroom",
                              humidity == "RH_6" ~ "north",
                              humidity == "RH_7" ~ "ironing",
                              humidity == "RH_8" ~ "teenager",
                              humidity == "RH_9" ~ "parents"))
```

As explained above, there are two variables that need this treatment; the temperature, and the humidity levels. In order

to plot the average monthly temperature in each room, I need to use `tidyr::pivot_longer()` (a little side note, I could have used `names_to = "room"`, instead of "temperature" and "humidity", but there's a reason for that. More on it below).

Now let's plot it:

```
flat_energy %>%
  mutate(month = month(date)) %>%
  group_by(month, temperature) %>%
  summarise(avg_temp = mean(temp_value)) %>%
  ggplot() +
  geom_line(aes(y = avg_temp, x = month, col = temperature)) +
  brottools::theme_blog()

## `summarise()` regrouping output by 'month' (override with `.groups` argument)
```



Ok great. But what if I had such a dataset per house for a whole city? How many datasets would that be? And how long would these operations take?

The first step I would take if I were in this situation, would be to write a function. I would make it general enough to work with temperature or humidity. Below is this function:

```
prepare_data <- function(energy, variable){

  variable <- enquo(variable)

  variable_label <- as_label(variable)
```

```

    regex_selector <- ifelse(variable_label == "temperature",
                             "T\\d{1}",
                             "RH_\\d{1}")
energy %>%
  pivot_longer(cols = matches(regex_selector),
               names_to = variable_label,
               values_to = paste0(variable_label, "_value")) %>%
  mutate(!!(variable) := case_when(grepl("1$", !!(variable)) ~ "kitchen",
                                     grepl("2$", !!(variable)) ~ "living",
                                     grepl("3$", !!(variable)) ~ "laundry",
                                     grepl("4$", !!(variable)) ~ "office",
                                     grepl("5$", !!(variable)) ~ "bathroom",
                                     grepl("6$", !!(variable)) ~ "outside",
                                     grepl("7$", !!(variable)) ~ "ironing",
                                     grepl("8$", !!(variable)) ~ "teenager",
                                     grepl("9$", !!(variable)) ~ "parents")) %>%

  mutate(month = month(date)) %>%
  group_by(month, !!(variable)) %>%
  summarise(across(.cols = ends_with("_value"),
                  .fns = mean),
            .groups = "drop")
}

```

This function does exactly the same thing as above:

```

prepare_data(energy, temperature) %>%
  ggplot() +
  geom_line(aes(y = temperature_value, x = month, col = temperature)) +
  brottools::theme_blog()

```



As you can see, I have the exact same plot. What's nice with this function, is that it uses many verbs from the `{tidyverse}` as well as the tidy eval framework for non-standard evaluation (which is why I did not use `names_to = "room"`, I wanted to use the variable label defined with `as_label()` and see if it works with `{tidytable}` as well).

Ok, so now let's imagine that I'm happy with this function, but I'd like it to run faster, and because I'm lazy, the less I have to modify it, the happier I am. This is where `{tidytable}` looks very promising. Let's rewrite the function to make it work with `{tidytable}`:

```

prepare_data_dt <- function(energy, variable){

  variable <- enquos(variable)

  variable_label <- as_label(variable)

  regex_selector <- ifelse(variable_label == "temperature",
                           "T\\d{1}",
                           "RH_\\d{1}")

energy %>%
  pivot_longer.(cols = matches(regex_selector),
                names_to = variable_label,
                values_to = paste0(variable_label, "_value")) %>%
  mutate.(!!(variable) := case_when(grepl("1$", !!(variable)) ~ "kitchen",
                                     grepl("2$", !!(variable)) ~ "living",
                                     grepl("3$", !!(variable)) ~ "laundry",
                                     grepl("4$", !!(variable)) ~ "office",
                                     grepl("5$", !!(variable)) ~ "bathroom",

```

```

      grepl("6$", !!(variable)) ~ "outside",
      grepl("7$", !!(variable)) ~ "ironing",
      grepl("8$", !!(variable)) ~ "teenager",
      grepl("9$", !!(variable)) ~ "parents")) %>%

mutate.(month = month(date)) %>%
summarise_across.(cols = ends_with("_value"),
                  .fns = mean,
                  .by = c(month, !!(variable))) %>%

ungroup()
}

```

As you can see, it's *almost* the same thing. `{tidytable}` verbs end with a `'.'` and that's it. Well almost (again), the biggest difference is how `{tidytable}` groups by a variable. It's very similar to how it's done in `{data.table}`, by using a `.by =` argument to verbs that support it, such as `summarise_across()` (which is also, by the way, another difference with standard `{tidyverse}` syntax). While I'll have to remember these, I'd argue that they're minor differences and if it can make my function run faster, I don't mind!

Now let's run a little benchmark. But first, let's define our data as a `tidytable` object:

```
energy_tidytable <- as_tidytable(energy)
```

Now we're good to go:

```

microbenchmark::microbenchmark(
  energy %>%
    prepare_data(temperature),
  energy_tidytable %>%
    prepare_data_dt(temperature),
  times = 10
)

## Unit: milliseconds
##              expr      min       lq      mean
##  energy %>% prepare_data(temperature)  847.9709  849.6671  868.6524
##  energy_tidytable %>% prepare_data_dt(temperature)  820.2051  838.6647  861.9685
##    median      uq      max neval
##  861.0652  880.8200  914.4685     10
##  858.9454  873.3268  936.0147     10

```

That is nice! It does indeed run faster, and with only some minor changes to the function! And how about using some more cores to run this function?

This can be done using `data.table::setDTthreads(n_cores)` where `n_cores` is the number of cores you want to use:

```

data.table::setDTthreads(12)
microbenchmark::microbenchmark(
  energy %>%
    prepare_data(temperature),
  energy_tidytable %>%
    prepare_data_dt(temperature),
  times = 10
)

## Unit: milliseconds
##              expr      min       lq      mean
##  energy %>% prepare_data(temperature)  832.9876  840.8000  874.3047
##  energy_tidytable %>% prepare_data_dt(temperature)  829.7937  831.2868  866.4383
##    median      uq      max neval
##  889.2684  898.6861  914.7178     10

```

```
## 836.8712 893.0613 997.8511 10
```

Maybe surprisingly, it did not run faster. It could very well be that my function does not really lend itself to running in parallel, and the overhead induced by distributing the work to the cpu cores cancels out the gains from running it in parallel. But in any case, this is really looking very interesting. I have not tested the whole package yet, but since the syntax is so similar to the `{tidyverse}`, you can try really quickly to see if the `{tidytable}` version of the function runs faster, and if yes, I don't really see a reason not to use it!